

CASTED: Core-Adaptive Software Transient Error Detection for Tightly-Coupled Cores

Konstantina Mitropoulou, Vasileios Porpodas,
Marcelo Cintra

University of Edinburgh

IPDPS 2013

Outline

- Transient Errors
- Software-based Error Detection
- Challenges/Existing approaches
- CASTED
- Performance & Fault Coverage Evaluation
- Conclusions

Transient Errors

As hardware errors become more frequent

↪ increased need for high-reliable and low-overhead error detection methodologies

Main sources of hardware errors :

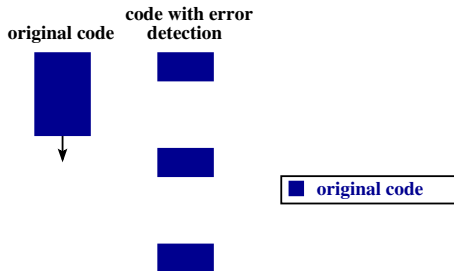
- small transistor technologies
- voltage scaling

Transient errors are:

- temporal phenomena
- the most frequent type of errors
- easy to handle at run-time

Compiler-based Error Detection (1)

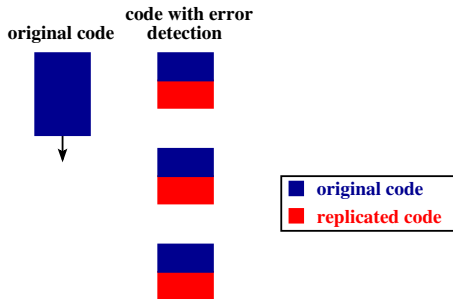
Dual-modular error
detection:



Compiler-based Error Detection (1)

Dual-modular error detection:

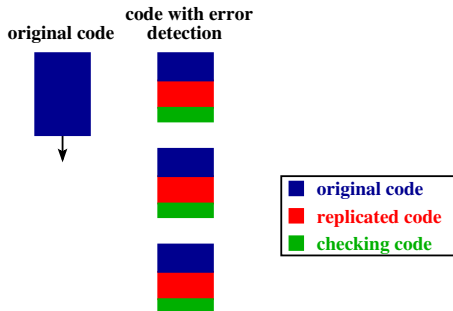
- *replicates* the computation



Compiler-based Error Detection (1)

Dual-modular error detection:

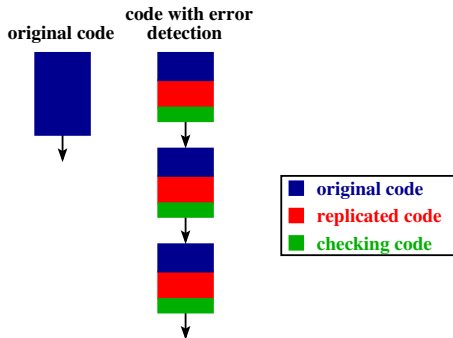
- *replicates* the computation
- *compares* the two outputs



Compiler-based Error Detection (1)

Dual-modular error detection:

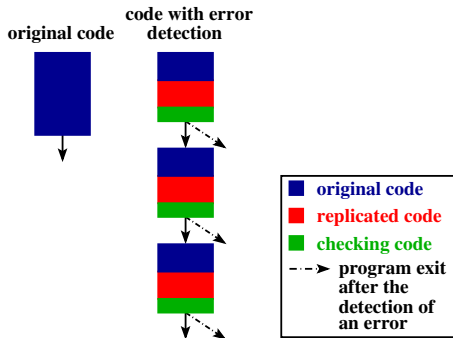
- *replicates* the computation
- *compares* the two outputs
- if the outputs are identical, then the execution *continues* normally



Compiler-based Error Detection (1)

Dual-modular error detection:

- *replicates* the computation
- *compares* the two outputs
- if the outputs are identical, then the execution *continues* normally
- in case of an error, the execution *rolls back* to the last checkpoint



Compiler-based Error Detection (2)

Challenge:

- decrease performance overhead without sacrificing system's reliability

Solutions/Existing approaches:

- optimize the code:
 - minimize checking points
 - reduce replicated code
- use more resources:
 - execute original and redundant code on separate cores

Overview of Existing Techniques

Overview of Existing Techniques



(a)

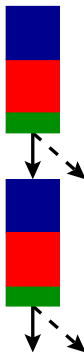
 **original code**

(a) Code without Error Detection (NOED)

Overview of Existing Techniques



(a)



(b)

■ original code
■ replicated code
■ checks
 - - ➔ program exit

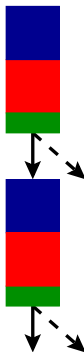
(a) Code without Error Detection (NOED)

(b) Single-Core Error Detection (SCED)

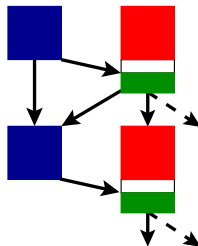
Overview of Existing Techniques



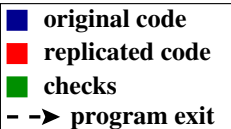
(a)



(b)

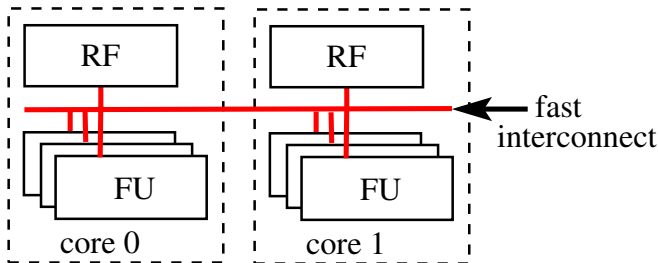


(c)



- | |
|---|
| (a) Code without Error Detection (NOED) |
| (b) Single-Core Error Detection (SCED) |
| (c) Dual-Core Error Detection (DCED) |

Limitations of Existing Techniques



Performance degradation factors:

- communication latency
- sub-optimal placement of the code
- lack of adaptivity

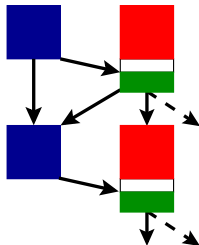
CASTED

CASTED solves this problem by introducing adaptation.

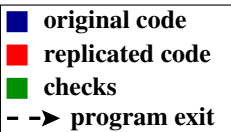
CASTED schedules the instructions taking into consideration:

- available resources
- communication latency

CASTED

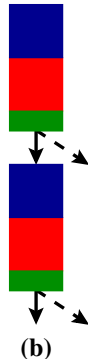
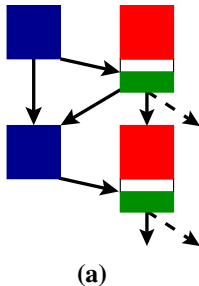


(a)



(a) Dual-Core Error Detection (DCED)

CASTED

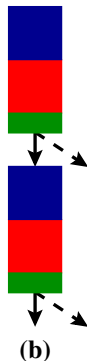
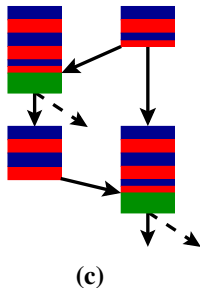
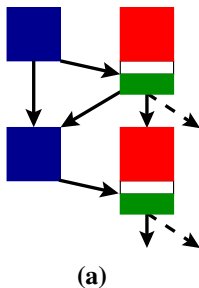


latency, resources

- original code
- replicated code
- checks
- -> program exit

- (a) Dual-Core Error Detection (DCED)
- (b) Single-Core Error Detection (SCED)

CASTED



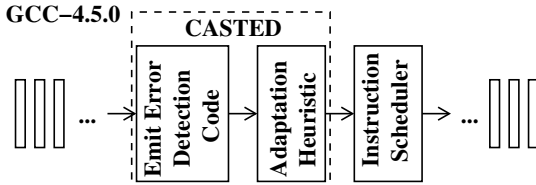
latency, resources

■ original code
 ■ replicated code
 ■ checks
 - -> program exit

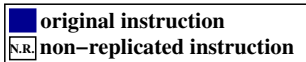
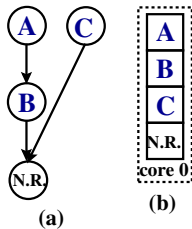
(a) Dual-Core Error Detection (DCED)
 (b) Single-Core Error Detection (SCED)
 (c) CASTED

CASTED Algorithm

- Emit error detection code:
 - *replicate* all necessary instructions
 - *isolate* original code from redundant code using register renaming
 - *insert* checks
- Adaptation heuristic:
 - is a greedy heuristic that maps the code to the current architecture configuration. It schedules the instructions considering the available resources, the communication latency and the data-flow.

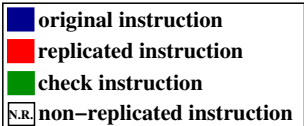
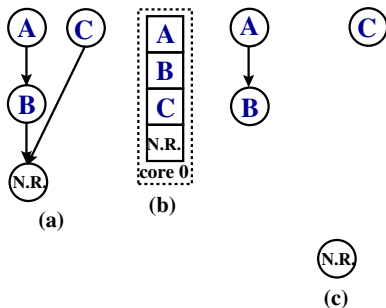


Example1 - Resource Constrained



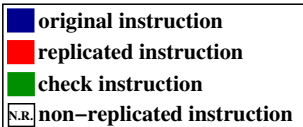
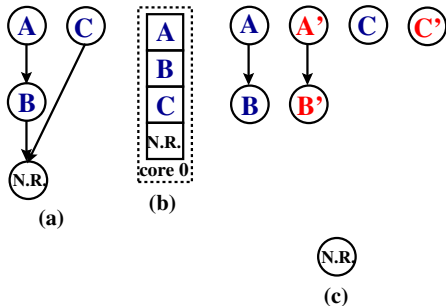
(a) Original Data Flow
 (b) Original Code without Error Detection (NOED)

Example1 - Resource Constrained



- (a) Original Data Flow
 (b) Original Code without Error Detection (NOED)
 (c) Data Flow with Error Detection Code

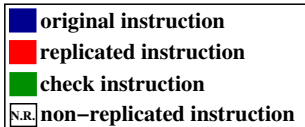
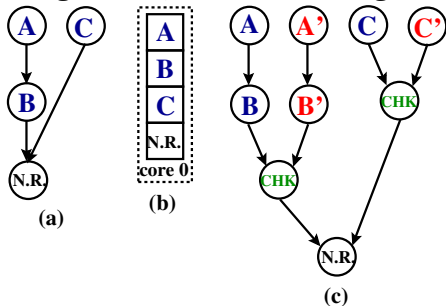
Example1 - Resource Constrained



- (a) Original Data Flow
 (b) Original Code without Error Detection (NOED)
 (c) Data Flow with Error Detection Code

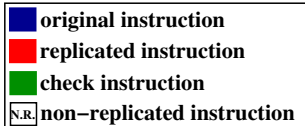
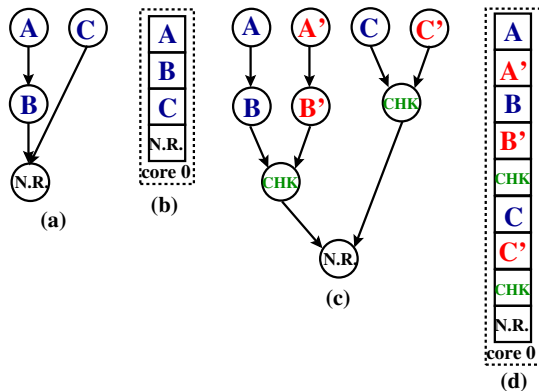
Example1 - Resource Constrained

Larger DFG, more ILP & longer critical path.



- (a) Original Data Flow
 (b) Original Code without Error Detection (NOED)
 (c) Data Flow with Error Detection Code

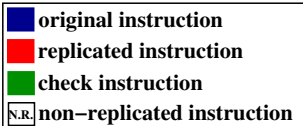
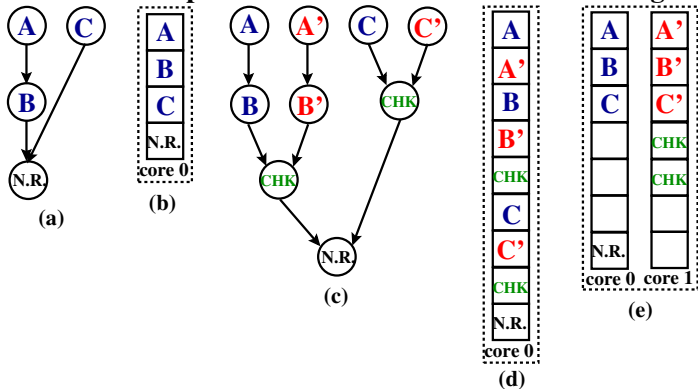
Example1 - Resource Constrained



- (a) Original Data Flow
 (b) Original Code without Error Detection (NOED)
 (c) Data Flow with Error Detection Code
 (d) Single-Core Error Detection (SCED)

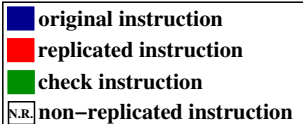
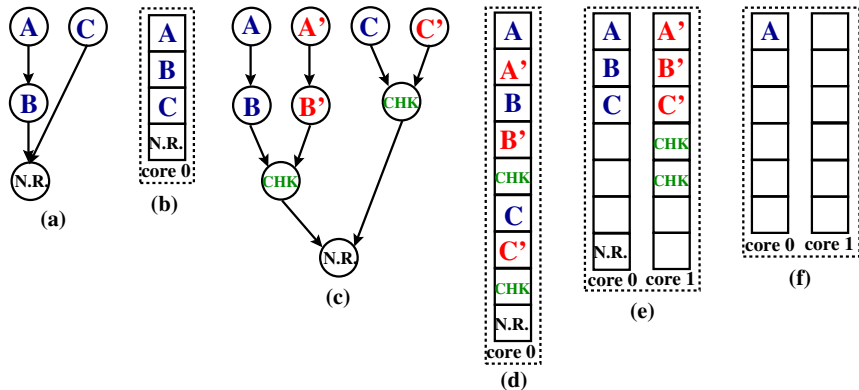
Example1 - Resource Constrained

Dual-core outperforms the resource constrained single-core.



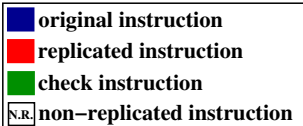
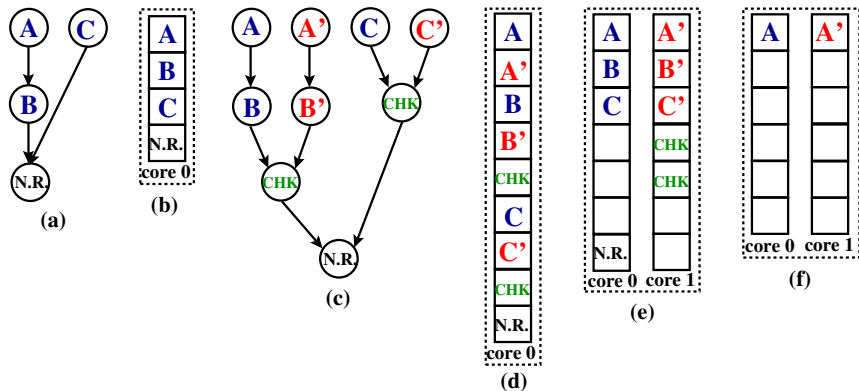
- (a) Original Data Flow
- (b) Original Code without Error Detection (NOED)
- (c) Data Flow with Error Detection Code
- (d) Single-Core Error Detection (SCED)
- (e) Dual-Core Error Detection (DCED)

Example1 - Resource Constrained



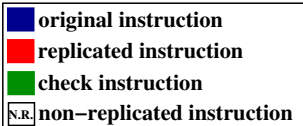
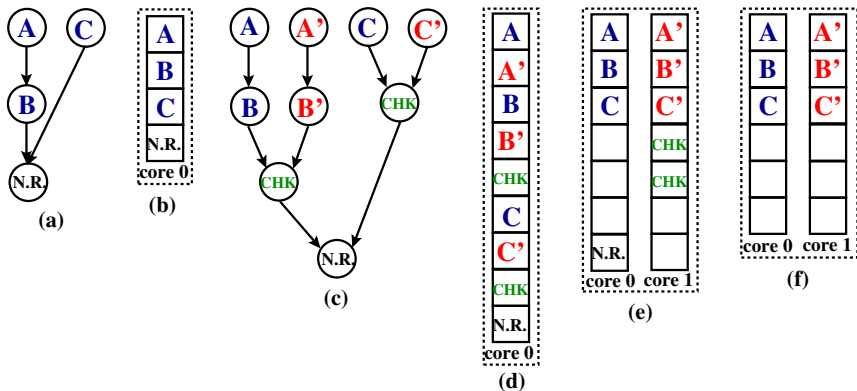
- (a) Original Data Flow
 (b) Original Code without Error Detection (NOED)
 (c) Data Flow with Error Detection Code
 (d) Single-Core Error Detection (SCED)
 (e) Dual-Core Error Detection (DCED)
 (f) CASTED

Example1 - Resource Constrained



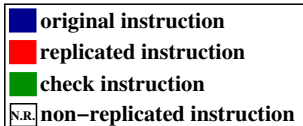
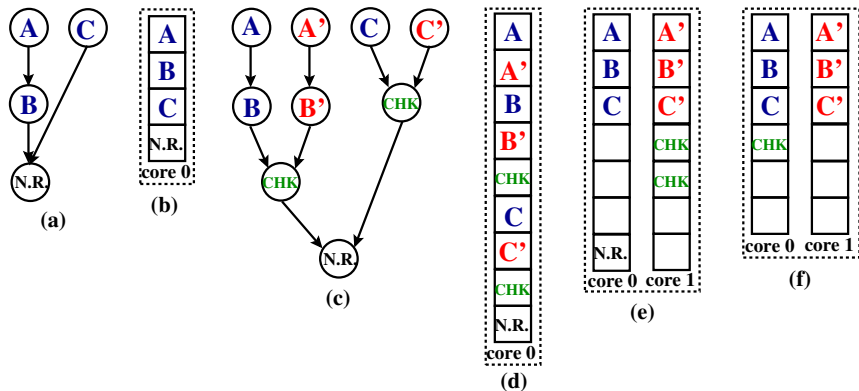
- (a) Original Data Flow
 (b) Original Code without Error Detection (NOED)
 (c) Data Flow with Error Detection Code
 (d) Single-Core Error Detection (SCED)
 (e) Dual-Core Error Detection (DCED)
 (f) CASTED

Example1 - Resource Constrained



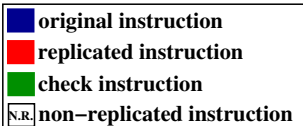
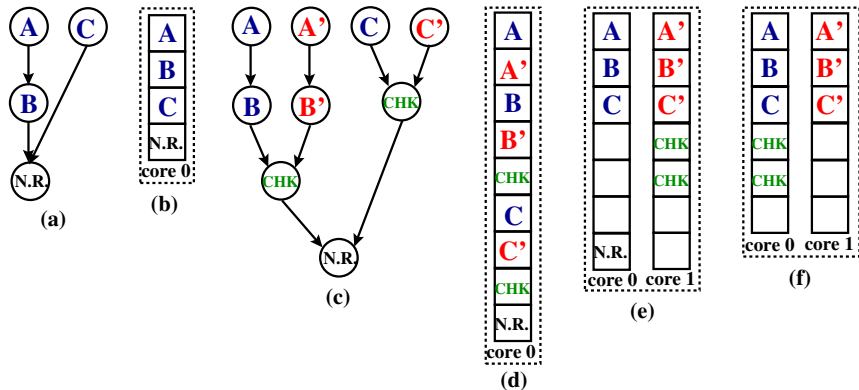
- (a) Original Data Flow
 (b) Original Code without Error Detection (NOED)
 (c) Data Flow with Error Detection Code
 (d) Single-Core Error Detection (SCED)
 (e) Dual-Core Error Detection (DCED)
 (f) CASTED

Example1 - Resource Constrained



- (a) Original Data Flow
 (b) Original Code without Error Detection (NOED)
 (c) Data Flow with Error Detection Code
 (d) Single-Core Error Detection (SCED)
 (e) Dual-Core Error Detection (DCED)
 (f) CASTED

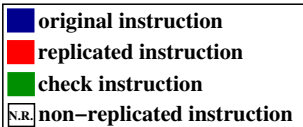
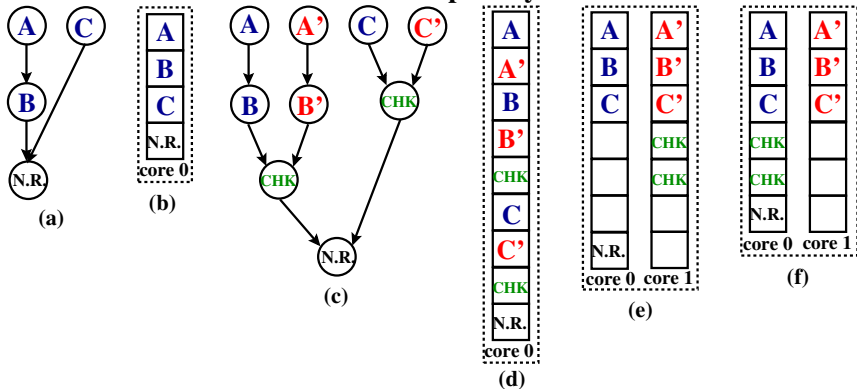
Example1 - Resource Constrained



- (a) Original Data Flow
 (b) Original Code without Error Detection (NOED)
 (c) Data Flow with Error Detection Code
 (d) Single-Core Error Detection (SCED)
 (e) Dual-Core Error Detection (DCED)
 (f) CASTED

Example1 - Resource Constrained

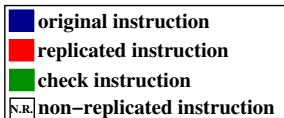
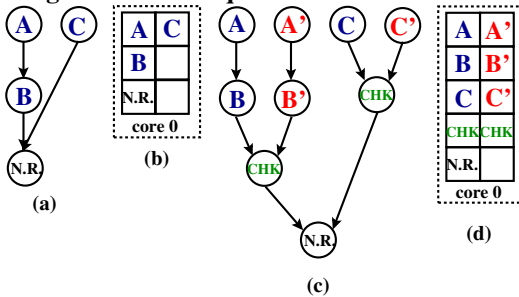
CASTED hides the communication penalty.



- (a) Original Data Flow
- (b) Original Code without Error Detection (NOED)
- (c) Data Flow with Error Detection Code
- (d) Single-Core Error Detection (SCED)
- (e) Dual-Core Error Detection (DCED)
- (f) CASTED

Example2 - Latency Constrained

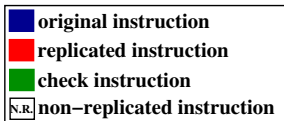
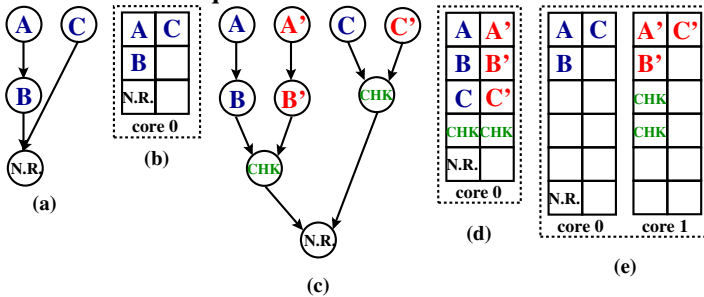
Single-core technique benefits from the extra resources.



- (a) Original Data Flow
 (b) Original Code without Error Detection (NOED)
 (c) Data Flow with Error Detection Code
 (d) Single-Core Error Detection (SCED)

Example2 - Latency Constrained

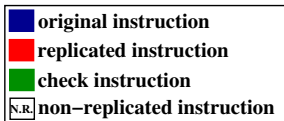
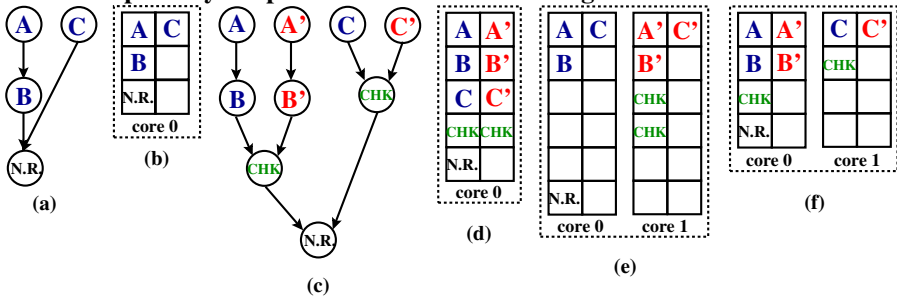
Dual-core technique benefits less from the extra resources.



- (a) Original Data Flow
- (b) Original Code without Error Detection (NOED)
- (c) Data Flow with Error Detection Code
- (d) Single-Core Error Detection (SCED)
- (e) Dual-Core Error Detection (DCED)

Example2 - Latency Constrained

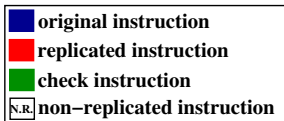
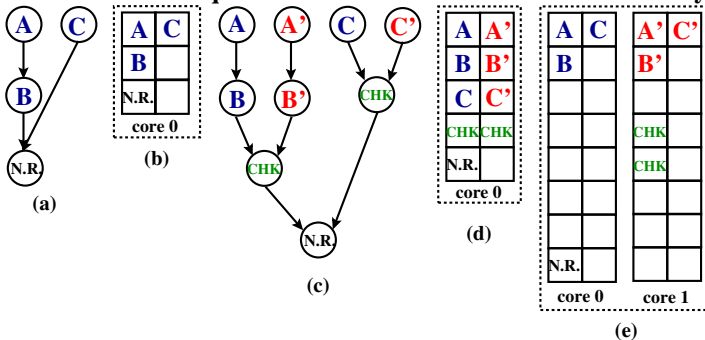
CASTED perfectly adapts to the architecture configurations.



- (a) Original Data Flow
 (b) Original Code without Error Detection (NOED)
 (c) Data Flow with Error Detection Code
 (d) Single-Core Error Detection (SCED)
 (e) Dual-Core Error Detection (DCED)
 (f) CASTED

Example2 - Latency Constrained

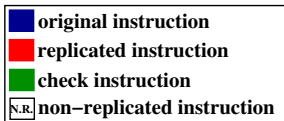
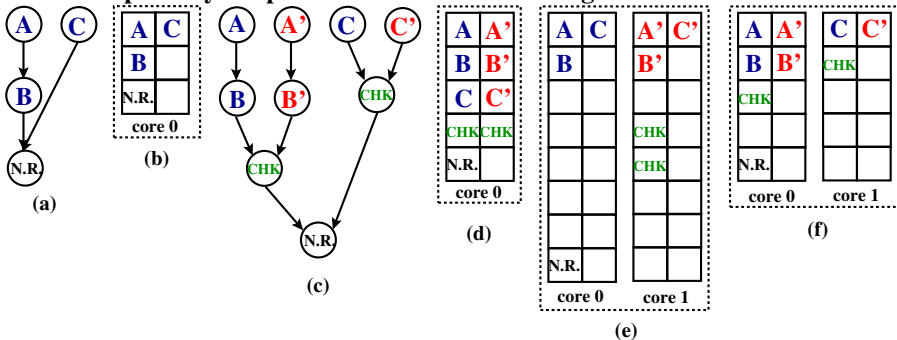
Dual-core technique suffers from the communication latency.



- (a) Original Data Flow
 (b) Original Code without Error Detection (NOED)
 (c) Data Flow with Error Detection Code
 (d) Single-Core Error Detection (SCED)
 (e) Dual-Core Error Detection (DCED)

Example2 - Latency Constrained

CASTED perfectly adapts to the architecture configurations.

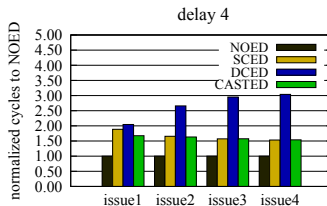
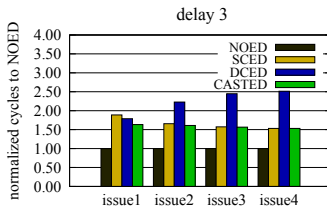
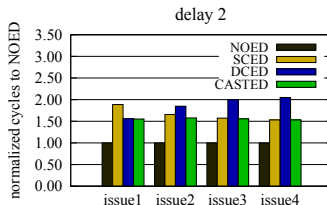
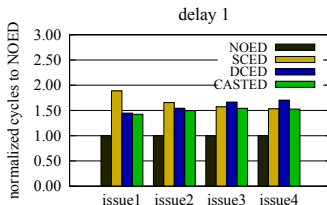


- (a) Original Data Flow
 (b) Original Code without Error Detection (NOED)
 (c) Data Flow with Error Detection Code
 (d) Single-Core Error Detection (SCED)
 (e) Dual-Core Error Detection (DCED)
 (f) CASTED

Experimental Set-up

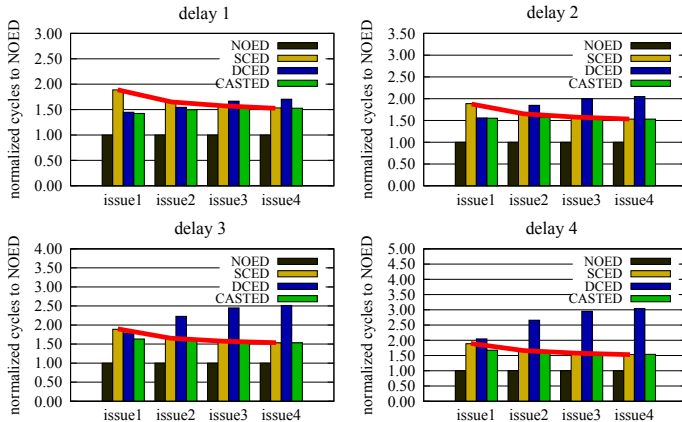
- Compiler
 - error detection code and adaptation passes added in the back-end of GCC-4.5.0
 - the *last* stage of CSE and DCE are turned-off
- Architecture
 - 2 IA64-based clusters whose issue-width takes values in the range of 1 to 4 and the communication latency varies from 1 to 4 cycles
 - SKI simulator
- Benchmarks
 - MediabenchII Video Benchmark suite
 - SPEC CINT2000
- Compare
 - NOED: No Error Detection (original code)
 - SCED: Single Core Error Detection
 - DCED: Dual Core Error Detection
 - CASTED: proposed technique

Performance Evaluation



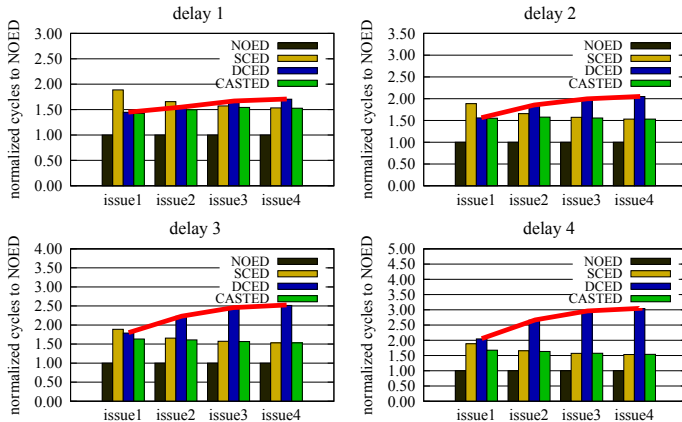
Performance Evaluation

SCED improves as the resources increase.



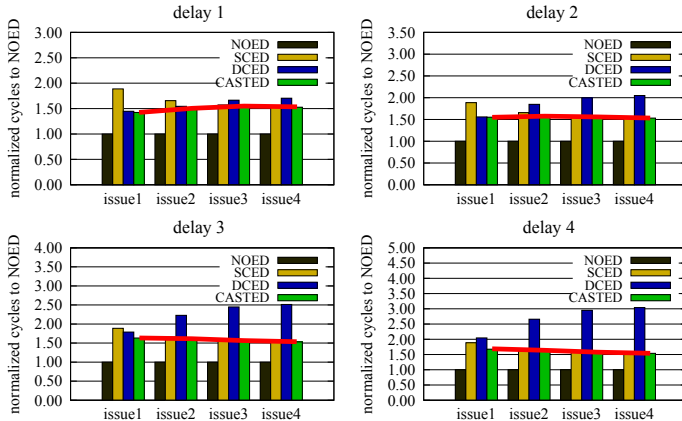
Performance Evaluation

DCED suffers communication latency and benefits less from the increase of issue-width.



Performance Evaluation

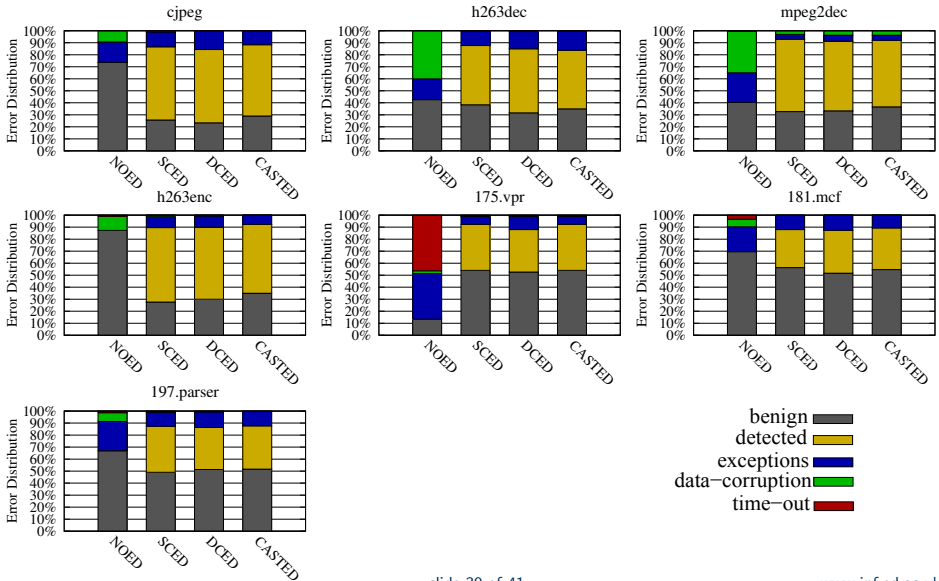
CASTED performs closely to the best technique for every configuration.



Fault Coverage Evaluation (1)

- Single-Event Upset (SEU) fault model
- Monte Carlo simulations:
 - 1 count dynamic instructions
 - 2 randomly pick one instruction
 - 3 randomly flip one bit of the instruction's output
 - 4 execute the program
 - 5 repeat steps 2-4 for 300 times for each implementation of each benchmark
- Errors taxonomy:
 - *benign errors*: result in correct output
 - *detected errors*: are the errors that a technique detects
 - *exceptions*: are the errors that raise exceptions
 - *data corrupt errors*: change program's output
 - *time-out errors*: result in infinite execution of the program.

Fault Coverage Evaluation (2)



Conclusions

- The overhead of the state-of-the-art techniques varies with the architecture configurations. More resources do not guarantee better performance.
- CASTED has a fixed overhead by optimally distributing the error detection overhead to the available resources.
- Performance tracks the best policy and sometimes outperforms it.
- No degradation in fault coverage.

CASTED: Core-Adaptive Software Transient Error Detection for Tightly-Coupled Cores

Konstantina Mitropoulou, Vasileios Porpodas,
Marcelo Cintra

University of Edinburgh

IPDPS 2013